

```
In [1]: import pandas as pd
```

```
In [2]: data=pd.read_csv("/home/placement/Downloads/TelecomCustomerChurn.csv")
```

```
In [3]: data.describe()
```

```
Out[3]:
```

| | SeniorCitizen | tenure | MonthlyCharges |
|-------|---------------|-------------|----------------|
| count | 7043.000000 | 7043.000000 | 7043.000000 |
| mean | 0.162147 | 32.371149 | 64.761692 |
| std | 0.368612 | 24.559481 | 30.090047 |
| min | 0.000000 | 0.000000 | 18.250000 |
| 25% | 0.000000 | 9.000000 | 35.500000 |
| 50% | 0.000000 | 29.000000 | 70.350000 |
| 75% | 0.000000 | 55.000000 | 89.850000 |
| max | 1.000000 | 72.000000 | 118.750000 |

In [4]: `data.head(10)`

Out[4]:

| | customerID | gender | SeniorCitizen | Partner | Dependents | tenure | PhoneService | MultipleLines | InternetService | OnlineSecurity | ... | DeviceProtec |
|---|------------|--------|---------------|---------|------------|--------|--------------|------------------|-----------------|----------------|-----|--------------|
| 0 | 7590-VHVEG | Female | 0 | Yes | No | 1 | No | No phone service | DSL | No | ... | |
| 1 | 5575-GNVDE | Male | 0 | No | No | 34 | Yes | No | DSL | Yes | ... | |
| 2 | 3668-QPYBK | Male | 0 | No | No | 2 | Yes | No | DSL | Yes | ... | |
| 3 | 7795-CFOCW | Male | 0 | No | No | 45 | No | No phone service | DSL | Yes | ... | |
| 4 | 9237-HQITU | Female | 0 | No | No | 2 | Yes | No | Fiber optic | No | ... | |
| 5 | 9305-CDSKC | Female | 0 | No | No | 8 | Yes | Yes | Fiber optic | No | ... | |
| 6 | 1452-KIOVK | Male | 0 | No | Yes | 22 | Yes | Yes | Fiber optic | No | ... | |
| 7 | 6713-OKOMC | Female | 0 | No | No | 10 | No | No phone service | DSL | Yes | ... | |
| 8 | 7892-POOKP | Female | 0 | Yes | No | 28 | Yes | Yes | Fiber optic | No | ... | |
| 9 | 6388-TABGU | Male | 0 | No | Yes | 62 | Yes | No | DSL | Yes | ... | |

10 rows × 21 columns



In [5]: data.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7043 entries, 0 to 7042
Data columns (total 21 columns):
#   Column                Non-Null Count  Dtype  
---  -
0   customerID            7043 non-null   object 
1   gender                 7043 non-null   object 
2   SeniorCitizen          7043 non-null   int64  
3   Partner                7043 non-null   object 
4   Dependents             7043 non-null   object 
5   tenure                 7043 non-null   int64  
6   PhoneService           7043 non-null   object 
7   MultipleLines          7043 non-null   object 
8   InternetService        7043 non-null   object 
9   OnlineSecurity         7043 non-null   object 
10  OnlineBackup           7043 non-null   object 
11  DeviceProtection       7043 non-null   object 
12  TechSupport            7043 non-null   object 
13  StreamingTV            7043 non-null   object 
14  StreamingMovies        7043 non-null   object 
15  Contract               7043 non-null   object 
16  PaperlessBilling       7043 non-null   object 
17  PaymentMethod          7043 non-null   object 
18  MonthlyCharges         7043 non-null   float64 
19  TotalCharges           7043 non-null   object 
20  Churn                  7043 non-null   object 
dtypes: float64(1), int64(2), object(18)
memory usage: 1.1+ MB
```

In [6]: data.shape

Out[6]: (7043, 21)

```
In [7]: list(data)
```

```
Out[7]: ['customerID',  
        'gender',  
        'SeniorCitizen',  
        'Partner',  
        'Dependents',  
        'tenure',  
        'PhoneService',  
        'MultipleLines',  
        'InternetService',  
        'OnlineSecurity',  
        'OnlineBackup',  
        'DeviceProtection',  
        'TechSupport',  
        'StreamingTV',  
        'StreamingMovies',  
        'Contract',  
        'PaperlessBilling',  
        'PaymentMethod',  
        'MonthlyCharges',  
        'TotalCharges',  
        'Churn']
```

```
In [8]: data["TotalCharges"] = pd.to_numeric(data['TotalCharges'], errors='coerce')
```

In [9]: data.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7043 entries, 0 to 7042
Data columns (total 21 columns):
#   Column                Non-Null Count  Dtype
---  -
0   customerID            7043 non-null   object
1   gender                 7043 non-null   object
2   SeniorCitizen          7043 non-null   int64
3   Partner                7043 non-null   object
4   Dependents             7043 non-null   object
5   tenure                 7043 non-null   int64
6   PhoneService           7043 non-null   object
7   MultipleLines           7043 non-null   object
8   InternetService        7043 non-null   object
9   OnlineSecurity          7043 non-null   object
10  OnlineBackup            7043 non-null   object
11  DeviceProtection        7043 non-null   object
12  TechSupport             7043 non-null   object
13  StreamingTV             7043 non-null   object
14  StreamingMovies         7043 non-null   object
15  Contract                7043 non-null   object
16  PaperlessBilling        7043 non-null   object
17  PaymentMethod           7043 non-null   object
18  MonthlyCharges          7043 non-null   float64
19  TotalCharges            7032 non-null   float64
20  Churn                   7043 non-null   object
dtypes: float64(2), int64(2), object(17)
memory usage: 1.1+ MB
```

```
In [10]: data1=data.drop(['customerID','StreamingTV','StreamingMovies','SeniorCitizen','Partner','Dependents','PhoneS
data1
```

```
Out[10]:
```

| | gender | tenure | MultipleLines | InternetService | OnlineBackup | TechSupport | Contract | MonthlyCharges | TotalCharges | Churn |
|------|--------|--------|------------------|-----------------|--------------|-------------|----------------|----------------|--------------|-------|
| 0 | Female | 1 | No phone service | DSL | Yes | No | Month-to-month | 29.85 | 29.85 | No |
| 1 | Male | 34 | No | DSL | No | No | One year | 56.95 | 1889.50 | No |
| 2 | Male | 2 | No | DSL | Yes | No | Month-to-month | 53.85 | 108.15 | Yes |
| 3 | Male | 45 | No phone service | DSL | No | Yes | One year | 42.30 | 1840.75 | No |
| 4 | Female | 2 | No | Fiber optic | No | No | Month-to-month | 70.70 | 151.65 | Yes |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 7038 | Male | 24 | Yes | DSL | No | Yes | One year | 84.80 | 1990.50 | No |
| 7039 | Female | 72 | Yes | Fiber optic | Yes | No | One year | 103.20 | 7362.90 | No |
| 7040 | Female | 11 | No phone service | DSL | No | No | Month-to-month | 29.60 | 346.45 | No |
| 7041 | Male | 4 | Yes | Fiber optic | No | No | Month-to-month | 74.40 | 306.60 | Yes |
| 7042 | Male | 66 | No | Fiber optic | No | Yes | Two year | 105.65 | 6844.50 | No |

7043 rows × 10 columns

```
In [18]: data2=data1.fillna(data1.median())
```

/tmp/ipykernel_5290/3414091449.py:1: FutureWarning: The default value of numeric_only in DataFrame.median is deprecated. In a future version, it will default to False. In addition, specifying 'numeric_only=None' is deprecated. Select only valid columns or specify the value of numeric_only to silence this warning.

```
data2=data1.fillna(data1.median())
```

```
In [11]: data2['Churn']=data2['Churn'].map({'Yes':1,'No':0})
```

In [19]: data2

Out[19]:

| | gender | tenure | MultipleLines | InternetService | OnlineBackup | TechSupport | Contract | MonthlyCharges | TotalCharges | Churn |
|------|--------|--------|------------------|-----------------|--------------|-------------|----------------|----------------|--------------|-------|
| 0 | Female | 1 | No phone service | DSL | Yes | No | Month-to-month | 29.85 | 29.85 | 0 |
| 1 | Male | 34 | No | DSL | No | No | One year | 56.95 | 1889.50 | 0 |
| 2 | Male | 2 | No | DSL | Yes | No | Month-to-month | 53.85 | 108.15 | 1 |
| 3 | Male | 45 | No phone service | DSL | No | Yes | One year | 42.30 | 1840.75 | 0 |
| 4 | Female | 2 | No | Fiber optic | No | No | Month-to-month | 70.70 | 151.65 | 1 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 7038 | Male | 24 | Yes | DSL | No | Yes | One year | 84.80 | 1990.50 | 0 |
| 7039 | Female | 72 | Yes | Fiber optic | Yes | No | One year | 103.20 | 7362.90 | 0 |
| 7040 | Female | 11 | No phone service | DSL | No | No | Month-to-month | 29.60 | 346.45 | 0 |
| 7041 | Male | 4 | Yes | Fiber optic | No | No | Month-to-month | 74.40 | 306.60 | 1 |
| 7042 | Male | 66 | No | Fiber optic | No | Yes | Two year | 105.65 | 6844.50 | 0 |


7043 rows × 10 columns

```
In [20]: data3=pd.get_dummies(data2)
data3
```

Out[20]:

| | tenure | MonthlyCharges | TotalCharges | Churn | gender_Female | gender_Male | MultipleLines_No | MultipleLines_No phone service | MultipleLines_Yes | Internet |
|------|--------|----------------|--------------|-------|---------------|-------------|------------------|-----------------------------------|-------------------|----------|
| 0 | 1 | 29.85 | 29.85 | 0 | 1 | 0 | 0 | 1 | 0 | |
| 1 | 34 | 56.95 | 1889.50 | 0 | 0 | 1 | 1 | 0 | 0 | |
| 2 | 2 | 53.85 | 108.15 | 1 | 0 | 1 | 1 | 0 | 0 | |
| 3 | 45 | 42.30 | 1840.75 | 0 | 0 | 1 | 0 | 1 | 0 | |
| 4 | 2 | 70.70 | 151.65 | 1 | 1 | 0 | 1 | 0 | 0 | |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 7038 | 24 | 84.80 | 1990.50 | 0 | 0 | 1 | 0 | 0 | 1 | |
| 7039 | 72 | 103.20 | 7362.90 | 0 | 1 | 0 | 0 | 0 | 1 | |
| 7040 | 11 | 29.60 | 346.45 | 0 | 1 | 0 | 0 | 1 | 0 | |
| 7041 | 4 | 74.40 | 306.60 | 1 | 0 | 1 | 0 | 0 | 1 | |
| 7042 | 66 | 105.65 | 6844.50 | 0 | 0 | 1 | 1 | 0 | 0 | |

7043 rows × 21 columns



```
In [21]: y=data3['Churn']
x=data3.drop('Churn',axis=1)
```


In [22]: y

```
Out[22]: 0      0
          1      0
          2      1
          3      0
          4      1
          ..
        7038     0
        7039     0
        7040     0
        7041     1
        7042     0
        Name: Churn, Length: 7043, dtype: int64
```

```
In [23]: from sklearn.model_selection import train_test_split
         x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.33,random_state=42)
```

```
In [24]: from sklearn.linear_model import LogisticRegression
         classifier=LogisticRegression()
         classifier.fit(x_train,y_train)
```

```
/home/placement/anaconda3/lib/python3.10/site-packages/sklearn/linear_model/_logistic.py:458: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html> (<https://scikit-learn.org/stable/modules/preprocessing.html>)

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression (https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result(
```

```
Out[24]: LogisticRegression()
```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [25]: y_pred=classifier.predict(x_test)
y_pred
```

```
Out[25]: array([1, 0, 0, ..., 1, 0, 0])
```

```
In [26]: from sklearn.metrics import confusion_matrix
confusion_matrix(y_test,y_pred)
```

```
Out[26]: array([[1518, 179],
               [ 271, 357]])
```

```
In [27]: from sklearn.metrics import accuracy_score
accuracy_score(y_test,y_pred)
```

```
Out[27]: 0.8064516129032258
```

```
In [ ]:
```